ChatGPT

AI Support for Neurodivergent Engineers in the Workplace

Executive summary

Neurodivergence is an umbrella term covering many “different styles of thinking, learning and communication,” including (but not limited to) ADHD, autism, dyslexia, dyscalculia, dyspraxia, and others; it is not a single condition and needs vary widely across individuals and contexts. Because you did not specify a particular neurodivergent condition, job family, or industry, this report focuses on patterns that recur across software/hardware engineering workplaces and highlights where needs are condition-specific or highly individualized.

Across research and practice, two themes matter most for engineering performance and wellbeing: (a) reducing **executive-function load** (planning, prioritizing, switching, remembering, initiating) and (b) reducing **communication/sensory friction** (ambiguity, meeting overload, interruptions, noise, and “tool sprawl”). Neurodivergent software professionals frequently report challenges in communication and collaboration, and also “excessive” tooling that drives distraction and cognitive overload.

AI can help substantially when it is deployed as *scaffolding* rather than as an ungoverned “autopilot”:

- For **junior engineers**, evidence increasingly shows the strongest gains from AI coding assistance and task scaffolding. A field experiment on developers found AI access increased completed tasks by about **26%** and that gains were **larger for less experienced developers**. A controlled experiment with GitHub Copilot found the treatment group completed a coding task **55.8% faster**, and the authors note heterogeneous effects suggesting AI pair programmers can help people “transition into software development careers.” In other job domains, a large organizational study also found generative AI improved productivity more for novices than for experienced workers (14% average, 34% for novices).
- For **senior engineers**, the highest leverage tends to be in: meeting/decision compression (summaries and action items), writing and rewriting (design docs, stakeholder updates), knowledge retrieval across systems, and delegation/mentoring workflows (turning plans into tickets, clarifying decisions, coaching juniors). Tools like intelligent meeting recap can generate AI meeting notes, recommended tasks, and topics/chapters, shifting senior attention from note-taking to decision-making and follow-up.

The same AI systems can also introduce material **limitations and risks**: hallucinated or insecure code, licensing/copyright ambiguity, data leakage, and new attack surfaces in agentic tooling. Research has documented Copilot recommending insecure code under certain conditions, and real-world analyses of AI-generated code have found substantial security issue prevalence in some languages and settings. Organizations should use a risk framework (for example, NIST’s AI Risk Management Framework and its Generative AI Profile) to define governance, privacy controls, and evaluation practices that are proportionate to the risk and context.

Managerially, the best evidence-based accommodations are not “AI-only.” They combine **work design** (clearer requirements, fewer interruptions, predictable schedules) with **access** to assistive and AI tools, and a psychologically safe culture where people can request adjustments without stigma. Recent survey research (CIPD) highlights that many neurodivergent employees do not disclose to their manager/HR due to privacy concerns, stereotypes, stigma, or career impact; this shapes how “opt-in” AI support and accommodations should be offered.

Junior and senior engineering workflows, tasks, and pain points

What “junior” and “senior” mean in practice

This section uses “junior” to mean engineers who are still building: domain fluency, codebase mental models, and confidence in ambiguity management; “senior” means engineers whose core value includes architecture judgment, cross-team coordination, and people/technical leadership. (Actual titles vary widely by company and industry.)

Junior engineers

Junior engineers often experience a high cognitive load from **unknown unknowns** (project context, standards, tooling, unwritten expectations) and from **rapid context switching** between learning, coding, and coordination. The “waiting for answers” and interruption pattern is especially important: a large survey-based analysis of professional engineers using Stack Overflow’s neurodiversity-related data reported that engineers with ADHD experienced more interruptions caused by waiting for answers. Qualitative work on neurodivergent professionals in software teams also reports that excessive tools and constant distractions can produce cognitive overload and degrade concentration.

Common junior role tasks and “AI-helpful” friction points are below.

Typical junior tasks - Interpreting a ticket/bug report; finding code ownership; reproducing issues. - Implementing scoped changes; writing tests; debugging failures in CI. - Reading internal docs and past PRs; following style and architecture constraints. - Communicating status and blockers; asking for help without overloading others.

Typical pain points where AI can help - Ambiguous requirements and unclear acceptance criteria (need structured questions). - Working memory load (“hold the state of the system in your head”) during debugging. - Fear of making mistakes in communication (status updates, PR descriptions). - Sustained attention challenges when forced into frequent micro-interruptions.

Junior workflow with AI scaffolding (illustrative)

```
flowchart TD
  A[Ticket / Request arrives] --> B[AI: summarize + extract acceptance criteria
+ unknowns]
  B --> C[Clarify questions to owner / stakeholder]
  C --> D[AI: propose 2-3 implementation approaches + risks]
  D --> E[Code + tests]
```

```
  E --> F[AI: generate test cases & edge cases]
  F --> G[Run CI / debug]
  G --> H[AI: explain error logs + suggest next debugging steps]
  H --> I[PR]
  I --> J[AI: draft PR description + change log + screenshots checklist]
  J --> K[Review feedback]
  K --> L[AI: turn feedback into actionable TODO list]
```

The key design principle for juniors is **“AI as a tutor + checklist, not a crutch.”** The strongest outcomes occur when AI accelerates learning loops and reduces the time spent stuck, while the engineer still practices reasoning and verification.

Senior engineers

Senior engineering work shifts from “do the task” to “shape the task system.” That includes architecture, technical strategy, high-quality reviews, and people leadership.

Typical senior tasks - Architecture/design authoring and review; making trade-offs; managing technical risk. - Mentoring and coaching; unblocking; improving team processes and documentation. - Cross-team alignment; running meetings; decision-making; stakeholder updates. - Incident leadership and postmortems; prioritization and roadmap shaping.

Typical pain points where AI can help - Meeting load and memory burden; converting discussion into decisions and action items. - Searching for information scattered across systems (docs, tickets, chats, repos). - Drafting high-quality written artifacts (design docs, strategy memos, ADRs). - Delegation overhead: turning goals into well-scoped tickets and coaching plans.

AI can reduce these costs by generating structured meeting recaps and action lists (for example, Teams intelligent recap provides AI notes and recommended tasks, plus chapters/topics and timeline markers). It can also help create and use “agents” to connect workflows to organizational knowledge (for example, Atlassian Rovo provides AI-powered search/chat/agents/studio connected through the Teamwork Graph, available to certain Jira/Confluence/JSM cloud plans).

Senior workflow with AI-assisted orchestration (illustrative)

```
flowchart LR
  A[Signals: incidents, metrics, customer asks, roadmap] --> B[AI: summarize +
cluster themes]
  B --> C[Decision memo: options, risks, dependencies]
  C --> D[Stakeholder alignment]
  D --> E[Delegation: break into epics/tickets + acceptance criteria]
  E --> F[Execution + reviews]
  F --> G[AI: summarize PRs + highlight risk areas for review]
  G --> H[Retrospective / postmortem]
```

```
H --> I[AI: draft action items + owners + due dates]
I --> A
```

For seniors, “AI advantage” is largely about **bandwidth reclamation**: fewer hours spent on transcription, status synthesis, and administrative writing, enabling more time for judgment, coaching, and systems thinking.

AI tools and features for neurodivergent engineers

Tool selection criteria that matter most for neurodivergent support

Because needs are heterogeneous, it helps to choose tools by *features that map to pain points* rather than by brand. The following are consistently high leverage across neurodivergent profiles:

- **Context control**: can the tool be scoped to a repo/project/workspace (reduces overwhelm and hallucination risk).
- **Explainability & grounding**: does it cite sources (tickets, docs) or provide links.
- **Interruption management**: does the tool reduce notifications or create structured summaries.
- **Privacy posture & retention controls**: critical for safe workplace use (see ethics section).
- **Workflow integration**: reduces context switching (not “one more tool”).

Comparison tables

Costs are shown only where they were clearly available in official vendor pages captured for this report; pricing changes frequently.

IDE and coding workflow assistants

Tool	Core features relevant to neurodivergent support	Pricing (official where available)	Best fit (junior vs senior)	Notes
GitHub Copilot	Code completions; chat/agent mode; CLI; higher tiers include “Copilot cloud agent” and “Copilot code review.”	Free: $0; Pro: $10/user/month; Pro+: $39/user/month.	Junior: **High**; Senior: **Med–High**	Strong for onboarding and unblocking; needs code review discipline due to security/license risks (see evidence).
Amazon Q Developer	IDE/CLI assistance; agentic requests; Pro tier expands limits at **$19/user/month**.	Free tier + Pro tier $19/user/month.	Junior: **High**; Senior: **Med**	Pricing page details tiering and agentic request limits; also offers pay-per-use for some transformations.

Tool	Core features relevant to neurodivergent support	Pricing (official where available)	Best fit (junior vs senior)	Notes
Tabnine	Emphasizes deployment flexibility (SaaS/VPC/on-prem/air-gapped) and "zero code retention"; adds agentic workflows at higher tier.	Code Assistant: $39/user/month (annual); Agentic Platform: $59/user/month (annual).	Junior: **Med**; Senior: **High** (esp. regulated environments)	Strong when privacy constraints limit public LLM tools.
Pair-programming process (AI-enabled)	AI as "rubber duck" + tutor: step-by-step debugging, test design, doc generation; reduces "stuck time."	Depends on model/tool	Junior: **High**; Senior: **High**	Empirically, AI coding tools can strongly reduce time-to-complete and help transitions into software dev careers.

Task, knowledge, and communication systems

Tool	What it does that matters for neurodivergent engineers	Pricing (official where available)	Best fit (junior vs senior)	Notes
Atlassian Rovo	AI-powered "search, chat, agents, and studio workflows" grounded in organizational knowledge via Teamwork Graph; available to certain Jira/Confluence/JSM cloud plans.	Included with qualifying cloud plans (Standard/Premium/Enterprise).	Junior: **Med**; Senior: **High**	Especially valuable for reducing "waiting for answers" by surfacing prior decisions and docs.
Notion (with AI features)	Unified docs/tasks/knowledge; business tier highlights AI agents + AI meeting notes; enterprise tier notes "zero data retention" by LLM providers for Notion AI.	Free: $0; Plus: $10/member/month; Business: $20/member/month; Enterprise: custom.	Junior: **Med**; Senior: **High**	Useful for externalizing working memory and creating structured personal workflows; privacy posture varies by plan.

Tool	What it does that matters for neurodivergent engineers	Pricing (official where available)	Best fit (junior vs senior)	Notes
Asana + Asana AI	Starter and Advanced plans include Asana AI; features include drafting tasks/status updates and building workflows; pricing $10.99/user/month annually for Starter, $24.99/user/month annually for Advanced.	Starter $10.99/user/month (annual); Advanced $24.99/user/month (annual).	Junior: **High** (task scaffolding); Senior: **High** (portfolio visibility)	Helpful for ADHD-like executive function friction; AI Studio adds credit-based automation.
Slack (AI features in plans)	AI channel/thread summaries; huddle notes; “Slack AI add-on” no longer separately purchasable; AI recap features available in paid plans (per Slack help).	Free $0; Pro and Business+ priced per user/month (varies; page shows examples such as $8.75/user/month for Pro).	Junior: **Med**; Senior: **High**	High value for reducing chat overload and catching up asynchronously.
Otter.ai	Meeting transcription + summaries; provides explicit plan pricing and minute limits.	Pro $8.33/user/month (annual); Business $24/user/month (annual) (plus other tiers).	Junior: **Med**; Senior: **High**	Particularly useful if your meeting platform lacks good native summaries.
Microsoft Teams Intelligent recap	AI notes + recommended tasks + topics/chapters; prerequisites include Teams Premium or Microsoft 365 Copilot licenses (audio recap requires Copilot).	Depends on Teams Premium / Copilot licensing	Junior: **High**; Senior: **High**	Converts meetings into actionable artifacts; reduces note-taking burden.
Google Meet “Take notes for me”	AI-powered note-taking; attaches notes doc to calendar event; licensing via Gemini plans/add-ons; notes follow Meet retention policy.	Via Gemini Enterprise / Education Premium / AI Meetings & Messaging add-on.	Junior: **High**; Senior: **High**	Strong for reducing working memory load and meeting fatigue.

Tool	What it does that matters for neurodivergent engineers	Pricing (official where available)	Best fit (junior vs senior)	Notes
Zoom AI Companion	Zoom positions AI Companion as available at no additional cost for paid accounts; newer “agentic” capabilities include scheduling follow-ups and generating documents.	Included with paid Zoom accounts (with constraints by region/plan).	Junior: **Med**; Senior: **High**	Can reduce follow-up burden and meeting admin load.
ChatGPT Business / Enterprise	Enterprise emphasizes encryption, custom retention, and “no training on your business data by default.”	Per-user per month; Enterprise is “contact sales.”	Junior: **High** (tutoring); Senior: **High** (writing/ decision support)	Useful “general-purpose” scaffold; governance and data controls are essential.

Accessibility, attention, and sensory supports

These tools matter even when “AI” is not branded as an LLM. For many neurodivergent engineers, accessibility features are core productivity infrastructure.

- **Windows Voice access** allows controlling a PC and authoring text with voice, explicitly without an internet connection.
- **Windows Live captions** provides captions across Windows 11 (including offline captions), and can include microphone audio to support in-person conversation comprehension; on Copilot+ PCs it can translate from many languages to English.
- **Windows Focus Sessions** are designed to minimize distractions and help keep tasks on track.
- **Apple Voice Control** supports navigating apps and dictating/editing text via voice.
- **Apple Live Captions** provides real-time transcription; Apple notes audio is processed on device, availability varies by language/region, accuracy varies, and should not be relied on in high-risk situations.
- Reading supports such as **Immersive Reader** and text-to-speech/speech-to-text can be valuable for dyslexia and sustained-attention challenges; Microsoft’s research summary cites work showing speech synthesis improved reading rates, comprehension, and ability to sustain attention for post-secondary students with dyslexia.

Evidence, case studies, measurable benefits, and limitations

What AI measurably improves

Coding throughput and task completion (strongest for juniors)
A field experiment (“Generative AI as a tool for…”) found developers with AI access completed **~26% more tasks**, with larger gains for less experienced developers and smaller/no gains for more experienced

developers; the same study reports a **~13.6% increase in commits**. This maps directly to junior pain points: “blocked time,” slow onboarding, and slower pattern recognition in unfamiliar codebases.

A controlled experiment on GitHub Copilot found a **55.8% faster** completion time for a specific coding task (HTTP server implementation), with the authors noting heterogeneous effects suggesting value in helping developers transition into software careers.

Learning and performance lift for lower-experience workers (cross-domain evidence)
In a large deployment in customer support, generative AI increased productivity by **14% on average**, with a **34% improvement for novice/low-skilled workers** and minimal impact for experienced/high-skilled workers. While not an engineering sample, it supports a consistent pattern: AI often acts as a *skill equalizer* for structured tasks, which is relevant to junior engineers and to neurodivergent engineers compensating for executive-function or communication friction.

Meeting attention, memory, and follow-up (high value for seniors and meeting-heavy teams)
Meeting summarization features aim to remove the need to simultaneously listen, decide, and write. Teams intelligent recap offers AI meeting notes and recommended tasks, plus topics/chapters and timeline markers. Google Meet’s “take notes for me” is explicitly positioned to let people focus on discussion while AI takes notes, attaching the notes document to the calendar event afterward. Zoom AI Companion similarly targets meeting summaries and follow-up automation.

These features are especially aligned with common neurodivergent pain points: working memory load, slower note-taking, auditory processing challenges, and fatigue from multi-hour meeting sequences.

What AI does not reliably improve, and where it can backfire

Security and quality risks in AI-generated code
Research has systematically evaluated when Copilot may recommend insecure code, raising the need for strong review and security tooling around AI-generated contributions. Empirical studies of AI-generated code in the wild have found substantial CWE instances and non-trivial vulnerability prevalence, varying by language and scenario. For neurodivergent engineers—especially those who may already struggle with time pressure, anxiety, or perfectionism—this can create a harmful dynamic: *AI increases speed, but the cost of verifying safety and correctness still exists* and may shift burden onto the engineer unless the team’s process absorbs it.

Licensing and attribution ambiguity
License compliance evaluation research finds even leading LLMs sometimes produce code “strikingly similar” to existing open-source implementations (reported ranges roughly 0.88%–2.01% in one benchmark), implying non-zero IP risk that must be managed.

Tool overload and cognitive overload
Neurodivergent professionals report that “excessive” tooling creates constant distractions and cognitive overload. Adding multiple AI layers (IDE assistant + chat + meeting bot + task bot) can worsen this unless the team deliberately reduces redundant tools and standardizes workflows.

Sensory and environmental constraints remain real
Workplace environment issues (especially auditory/social stimuli) can reduce job satisfaction and

performance among autistic employees, with perceived environmental control acting as a key mechanism. AI does not fix noise, poor lighting, or open-office interruption norms; it can only partially compensate (e.g., captions, better async documentation).

Role-specific evidence on neurodivergent engineers' lived experience

A Microsoft Research study of neurodiverse software engineering employees combined interviews with neurodiverse technologists and a survey of **846** engineers (59 identifying as neurodiverse). It highlights that many employees had not disclosed neurodiversity status, and it points to workflow/practice changes that could better support participation (including improving communication tooling and processes). Recent qualitative research with neurodivergent professionals in software development teams also identifies communication difficulties, prejudice, and tool overload as key challenges.

Survey research (CIPD) suggests disclosure barriers remain common and are closely tied to fears of stigma and stereotypes—directly affecting how teams should roll out AI supports (opt-in, privacy-preserving, not forcing disclosure).

Does neurodivergence become an advantage or disadvantage in the AI era?

The most defensible answer is: **both**, depending on the environment and on which parts of the job AI changes.

Neurodivergence can confer advantages in contexts where the role rewards deep focus, pattern recognition, and novel ideation. Policy and business-oriented reports (for example, the UK Buckland Review) note autistic strengths that may outperform non-autistic peers in certain contexts, including attention to detail and focus. Academic synthesis also cautions against simplistic "autism advantage" claims: a critical/systematic review evaluates the evidence base for superior workplace performance and indicates it is not uniformly established. Similarly, research on ADHD and creativity is nuanced: one study finds higher ADHD symptoms in the general population associated with stronger divergent thinking measures, but that does not automatically translate to better job outcomes without supportive work design. For dyslexia, a review on "spatial advantage" suggests little evidence for a generalized spatial reasoning advantage and often equal or worse performance on such tasks—an important corrective to stereotypes.

In the AI era, the relative advantage/disadvantage shifts because AI changes which skills are scarce:

- **AI reduces the penalty of communication-heavy norms**. Writing assistants, summarizers, and structured prompt templates can reduce the "communication tax" neurodivergent engineers often pay (especially when instructions are ambiguous or meetings are unstructured).
- **AI increases the importance of verification, judgment, and risk management.** If an engineer's strengths include careful analysis and attention to detail, they may be well-positioned—*but only if time is allocated for review and quality*.
- **AI can magnify disadvantages in surveillance-heavy cultures.** If productivity measurement becomes "AI output per day" rather than impact, neurodivergent engineers who work in bursts or rely on recovery time can be penalized. CIPD's findings about stigma and reluctance to disclose underline why organizations must protect psychological safety and evaluate fairly.

- **AI’s strongest measured benefits often accrue to novices.** This can narrow the early-career gap for neurodivergent juniors—*if* the team supports safe learning and does not punish “asking for help,” which AI may partially replace.

Overall: neurodivergent engineers are likely to do better relative to peers in the AI era **when** AI reduces irrelevant friction (admin, transcription, boilerplate) and the workplace invests in neuroinclusive practices. When AI is introduced as uncontrolled automation plus heightened monitoring, it can worsen overload and inequity.

Accommodations and best practices for managers and teams

What good neuroinclusive practice looks like with AI in the loop

Survey evidence (CIPD) indicates organizations that take action toward neuroinclusion more often report positive impacts on employee wellbeing, comfort discussing neurodiversity, appreciation for different thinking styles, and other outcomes. The same report highlights specific practices employers see as having strong positive impact, including **flexible working**, **clear access to reasonable adjustments**, and **education/awareness raising**.

A practical implication: AI tools should be treated as one component of “reasonable adjustments,” not as a replacement for them.

Concrete practices that reduce friction for both junior and senior engineers

Use these as defaults; tailor via discussion rather than requiring disclosure of diagnosis.

Make work more explicit - Requirements: define “done,” include examples, edge cases, and acceptance tests (reduces ambiguity-driven stress and rework). This directly addresses common neurodivergent challenges around unclear instructions and unstructured communication reported in software team studies.
- Decisions: record decisions asynchronously (ADR / decision memo) and link them in tickets and PRs; AI can draft these, but people must validate.

Standardize team communication - Use meeting notes automation by default when allowed (Teams recap, Meet notes) and always capture action items/owners.
- Prefer async updates for status; use AI to turn raw notes into structured updates to reduce communication burden.

Reduce interruption load - Define “office hours” + async channels; explicitly protect focus blocks. ADHD-related survey findings about interruptions (e.g., waiting for answers) suggest this is not a minor issue.
- Encourage use of system-level focus tools (e.g., Focus Sessions) to suppress notifications during deep work.

Engineer the environment - Sensory control options (quiet zones, remote days, noise-canceling allowances, predictable desk setups) matter because sensory and social stimuli can be particularly problematic and correlate with lower satisfaction/performance.

Make career progression legible - Provide explicit expectations for each level; ensure feedback is concrete and behaviorally anchored, not "vibes-based."

How managers should roll out AI supports without forcing disclosure

CIPD reports neurodiversity is talked about by managers in only **31%** of organizations, and many neurodivergent employees have not told their manager/HR due to privacy, stigma, stereotypes, and career impact concerns. This implies:

- Offer AI supports as **universal** productivity tooling ("available to everyone"), and let individuals opt-in to the parts that help.
- Build an **accommodations menu** that doesn't require diagnosis: captions, meeting summaries, flexible hours, written agendas, fewer meetings, async-first norms.
- Ensure **psychological safety**: do not use AI adoption data as a proxy for performance, and do not require employees to justify tool usage.

Ethical, privacy, bias, and security considerations

Privacy and confidentiality

Workplace AI often touches sensitive material (source code, security vulnerabilities, customer data, HR discussions). Tools differ in retention and training policies:

- Notion notes that Notion AI does not use customer data to train models unless the customer agrees to share data; its Enterprise plan states LLM providers follow a "zero data retention" policy for Notion AI.
- OpenAI's business pricing page for ChatGPT Enterprise highlights "encryption at rest and in transit," custom retention policies, and "no training on your business data by default."
- Apple's Live Captions documentation states audio is processed on device, and it warns accuracy varies and should not be relied on in high-risk situations—important for workplace expectations.
- Windows features like Voice access and Live captions emphasize offline capability, which can be privacy-beneficial in some scenarios.

Best-practice privacy pattern: treat AI prompts/outputs as potentially discoverable corporate records; prohibit pasting secrets into tools without explicit approval; and prefer enterprise-managed offerings with clear retention controls.

Bias, fairness, and evaluation

Bias risks appear in two places:

- **Model outputs**: language tone suggestions, "professionalism" scoring, and summarization can encode cultural bias (what counts as "clear," "polite," "executive-ready").
- **Workplace measurement**: if AI makes output more verbose or faster, managers may misinterpret speed as competence or penalize those who use accessibility tools.

A risk management approach like NIST's AI RMF and its Generative AI Profile helps organizations connect governance actions to concrete risks (privacy, security, harmful bias) rather than ad hoc rules.

Secure use of AI code generation and agents

Given research on insecure AI-generated code and licensing ambiguity, teams should adopt “secure-by-default” workflow controls:

- Treat AI-generated code as “untrusted” until reviewed and tested; security research shows AI suggestions can include insecure patterns.
- Add automated security scanning (SAST, dependency scanning, secret scanning) and require human code review, especially for authentication, crypto, and serialization code paths. (This complements research findings of non-trivial CWEs in AI-generated code in repositories.)
- Add IP compliance checks where relevant; license compliance evaluation work suggests a measurable (even if relatively low) rate of near-duplicate code generation.

Practical action checklist and prioritized next steps

Checklist for neurodivergent engineers

Keep this lightweight and iterative—optimize for *reducing cognitive load* rather than “using more tools.”

- Choose **one** core assistant for coding (Copilot / Q Developer / Tabnine) and **one** core system for externalizing tasks and notes (Asana / Notion / Jira+Rovo), to avoid tool overload.
- Build a personal “prompt pack”:
- “Summarize this ticket into acceptance criteria + unknowns.”
- “Turn this review feedback into a prioritized checklist.”
- “Explain this error log and propose the next 5 debugging steps.”
- Use meeting capture by default (if policy allows): enable recap/notes so you can focus on comprehension and participation rather than transcription.
- Use OS-level supports as needed:
- Captions for comprehension (Windows Live captions / Apple Live Captions).
- Voice access/control for low-friction text entry (Windows Voice access / Apple Voice Control).
- Focus blocks (Windows Focus Sessions or equivalent) during deep work.
- Add a “verification ritual” for AI code:
- Run tests + linters, review for security-sensitive patterns, and ask the AI to generate edge cases—but always validate.

Checklist for managers and teams

- Make “clarity artifacts” mandatory: written acceptance criteria, decision notes, and post-meeting action lists (AI can draft; humans own correctness).
- Default to **async-first**: record meeting summaries and allow written participation; this reduces the penalty of real-time processing differences.
- Reduce interruptions structurally: defined support hours, bounded response-time expectations, and protected focus time.
- Offer accommodations without requiring disclosure: captions, flexible schedules, quieter workspaces, and accessible documentation are universally useful.
- Implement AI governance: approved tools list, data-handling guidance, and secure SDLC policies aligned to a risk framework (e.g., NIST GenAI Profile).

Prioritized next steps

First priority: standardize “low-risk, high-return” supports - Turn on meeting notes/summaries where licensed and policy-approved (Teams intelligent recap, Google Meet notes, Zoom AI Companion) to reduce cognitive load and improve follow-through.
- Adopt OS-level accessibility defaults (captions, dictation/voice control availability) as normal productivity options, not exceptional accommodations.

Second priority: deploy coding assistants with guardrails - Pilot an IDE assistant (Copilot / Amazon Q Developer / Tabnine) with explicit rules: no secrets in prompts, mandatory review, security scanning, and templates for safe usage.
- Measure outcomes that matter: cycle time, reopened bugs, review load, and self-reported cognitive load—because speed alone can be misleading.

Third priority: reduce tool overload and strengthen knowledge systems - Consolidate sources of truth and deploy organization-grounded search/chat (e.g., Rovo-style knowledge connection) to cut “waiting for answers” interruptions.
- Create reusable onboarding packs (AI-generated drafts reviewed by seniors) to speed junior ramp and reduce ad hoc mentorship load.

Fourth priority: institutionalize neuroinclusive culture - Train managers and teams (awareness + concrete practices). CIPD data suggests awareness is often lacking and that many employees avoid disclosure due to stigma and career concerns—this is a structural barrier, not an individual failing.
