

AI-Mediated Knowledge and the New Office Politics in Large Software Companies

Executive summary

AI systems that “work” only when an organization’s knowledge is written down, indexed, permissioned, and retrievable (enterprise search, RAG-enabled LLM assistants, knowledge graphs, agent toolchains) reshape office politics because they convert **informal, person-held expertise** into **legible, queryable, and auditable** organizational capital. This shifts influence away from people who previously controlled “tribal knowledge” and toward those who control the **knowledge supply chain**: authoring and “source of truth” documents, taxonomy/metadata, ingestion/connectors, retrieval/ranking, access controls, and audit/compliance. The effect is not simply “more efficient knowledge work”; it is a redistribution of informational power and a reconfiguration of turf boundaries. ¹

Two empirical patterns matter for politics. First, field evidence suggests generative AI assistance can **raise productivity with large heterogeneity**, often benefiting less-experienced workers more than experts—compressing the performance advantage that senior specialists traditionally used as a base of power. In a 5,179-agent deployment study, access to a generative AI assistant increased issues resolved per hour by **14% on average**, including **34% for novice/low-skilled workers**, with “minimal impact” on experienced/high-skilled workers. ² Second, controlled field experiments with high-skill professionals show a “*jagged frontier*”: for tasks inside an AI capability frontier, AI increased speed (~**25%**), quality (>**40%**), and task completion (~**12%**), with larger gains for below-average performers (**43%**) than above-average performers (**17%**); yet outside the frontier, overreliance can degrade correctness. ³

The politics emerge because enterprise AI is not a monolith; it is a pipeline. Vendor documentation makes explicit that enterprise assistants are typically **permission-aware** (they “only surface what you can access”), and they often store interaction histories (prompts/responses) and provide administrative controls for retention and compliance—creating a new visibility layer and “audit culture” around knowledge work. ⁴

Practical advantage in this new regime comes from becoming (or aligning with) at least one of these roles: **knowledge product owner** (canonical docs), **index/connector owner** (what gets ingested), **retrieval/ranking tuner** (what gets privileged), **access-control steward** (who can see what), or **audit/metrics interpreter** (what leaders believe is happening). These are power levers in French & Raven’s sense of informational/expert/legitimate power—now partly mediated by algorithms. ⁵

This report provides an analytical model of these shifts, the organizational attributes that determine who wins which turf wars, a tactics-and-risks matrix (including overt and covert tactics), mitigations and long-term career implications, and a phased roadmap for adoption.

How AI that depends on written knowledge reshapes power

Enterprise LLM assistants increasingly operate as **retrieval-augmented** systems: they combine a generative model with a retriever over a document index (dense vectors, keyword search, or knowledge graphs). Retrieval-augmented generation (RAG) was formalized as a model class that pairs parametric memory (the LLM) with non-parametric memory (a retrievable corpus), explicitly motivated by limits in precision, updating, and provenance. ⁶ In practice, vendors describe the same pattern: Microsoft 365 Copilot coordinates LLMs with content in Microsoft Graph that the user has permission to access. ⁷

Because the model's usefulness depends on what is written and retrievable, organizations become structurally incentivized to (a) document more, (b) standardize and index, and (c) govern access and retention. Microsoft's Copilot readiness guidance explicitly ties answer quality and security to data being "well governed, current, and appropriately shared." ⁸ This is a technical statement with political consequences: whoever determines "governed/current/appropriately shared" effectively influences what the AI claims is true inside the company.

The new "knowledge supply chain" and its control points

The core office-politics innovation is that decision-making becomes increasingly mediated by systems that **select what is relevant**. Scholarship on algorithms emphasizes that algorithms (including search/ranking) are not neutral—they adjudicate relevance and visibility, shaping what is encountered and acted upon. ⁹ This is amplified when LLMs summarize across sources: the "winner" is not merely whose document is found, but whose framing becomes the AI's synthesized narrative.

Kellogg, Valentine, and Christin describe algorithmic technologies in organizations as a "contested terrain of control," where new control mechanisms invite adaptation and resistance by workers. ¹⁰ In knowledge settings, the analogous contested terrain is the end-to-end pipeline:

- **Authoring power:** who can create and bless canonical artifacts (design docs, runbooks, policies).
- **Curation power:** who controls taxonomy, metadata, templates, and what counts as "done."
- **Ingestion power:** who runs connectors, decides what gets indexed, and how frequently.
- **Retrieval/ranking power:** who can tune ranking signals, embeddings, freshness, and "featured sources."
- **Access-control power:** who sets ACLs/ABAC rules and prevents—or enables—oversharing.
- **Audit/analytics power:** who can see usage, prompts, admin actions, and compliance logs—and how they interpret them.

These map cleanly onto classical "bases of power" (expert/informational/legitimate/reward), but are now partly mediated by tools and policies rather than purely interpersonal networks. ¹¹

Why performance heterogeneity changes status hierarchies

Generative AI assistance can erode a traditional status hierarchy: "senior people know more and thus are indispensable." Brynjolfsson, Li, and Raymond's field study reports that the largest productivity gains accrue to novice/low-skill workers (34%) relative to the average (14%), with minimal effect for experienced/high-skilled workers, suggesting AI "disseminates best practices" and speeds up the experience curve. ²

Dell'Acqua and colleagues similarly find that AI especially benefits lower performers (43% vs 17% for higher performers) on tasks within the capability frontier. ¹²

This redistribution has political implications:

- **Experts lose monopoly rents** on routine know-how if that know-how is encoded in accessible artifacts and the AI can surface it well.
- **Newcomers gain leverage** because they can operate with higher baseline competence, ask better questions, and produce higher-quality drafts faster.
- **Gatekeeping shifts upstream:** experts who adapt can preserve influence by becoming doc owners, reviewers, evaluators, or governance stewards, especially where tasks are on the “jagged frontier” and require judgment. ¹³

Diagram of information flow and power shift

```
flowchart LR
    subgraph PreAI["Pre-AI knowledge regime (people-centric)"]
        A["Tribal knowledge\nin experts' heads"] --> B["Social networks\n& meetings"]
        B --> C["Decisions & execution"]
        A -. gatekeeping .-> C
    end

    subgraph AIRegime["AI-mediated knowledge regime (artifact-centric)"]
        D["Written artifacts:\nrunbooks, design docs,\ntickets, chats"] --> E["Ingestion:\nconnectors + ETL"]
        E --> F["Index:\nkeyword + vector + KG"]
        F --> G["Retrieval/ranking"]
        G --> H["LLM assistant\n(RAG synthesis)"]
        H --> I["Decisions & execution"]
    end

    subgraph ControlPoints["New control points (power levers)"]
        CP1["Doc ownership\n& templates"]
        CP2["Connector scope\n& cadence"]
        CP3["Ranking signals\n& featured sources"]
        CP4["Access controls\n(RBAC/ABAC)"]
        CP5["Audit logs\n& compliance analytics"]
    end

    CP1 --> D
    CP2 --> E
    CP3 --> G
    CP4 --> H
    CP5 --> I
```

The politics arise because each control point can be contested: teams fight over who owns templates, who runs connectors, which sources are “authoritative,” and who can see which audit artifacts.

Attributes that determine turf wars and who gains influence

This section enumerates the key attributes you requested and analyzes how each shapes internal power dynamics.

Organizational structure and knowledge topology

The primary structural variable is whether “knowledge” is treated as:

- **Centralized:** a central team owns the primary platform (wiki/search/LLM), sets templates, controls ingestion and ranking, and often becomes the default arbiter of “truth.”
- **Federated:** domain teams own their knowledge as “products,” while a smaller platform team provides self-serve infrastructure; governance is federated and rule-based across domains (rather than one group manually policing everything). This mirrors the “data mesh” framing: domain-oriented ownership plus federated computational governance. ¹⁴

A useful way to predict turf wars is to treat this as a **decision-rights problem**. Data governance research frames governance as identifying decision domains and assigning who decides. ¹⁵ When applied to AI-mediated knowledge, the decision domains often become:

- What artifacts count as authoritative?
- Who can publish/approve?
- What gets indexed, and by whom?
- What ranking/grounding policies apply?
- What retention and deletion rules apply?
- Who audits, and with what thresholds?

If these decision rights are ambiguous, politics rushes in. Knowledge-management scholarship explicitly warns that knowledge flows are embedded in power relations, and power dynamics can be both a barrier and an enabler. ¹⁶

Types of AI systems and their political “attack surfaces”

Different AI system designs create different “power surfaces” (places where influence accumulates).

Knowledge bases and wikis (human-authored, curated) Enterprise wiki research shows these systems split participation into distinct behaviors (knowledge creation vs integration), with different motivating factors and organizational levers. ¹⁷ Politically, wikis create “documentation real estate” that can be captured: owning the main page, the canonical runbook, the template library.

Enterprise search and ranking Search is a gatekeeping mechanism: relevance ranking influences what is seen, remembered, and acted upon, and is therefore a power lever. ¹⁸ Internal “SEO” behaviors (keyword stuffing, carefully chosen titles, freshness gaming) become office-politics tactics.

LLM assistants (RAG, agents, tool-using) Copilot-style systems explicitly connect LLMs to enterprise content and context, usually through a graph and connectors. ¹⁹ These systems introduce new politics: whose documents get cited, which agent is invoked, and what the assistant recommends becomes a quasi-managerial influence channel.

Access controls and permission-aware retrieval Most serious enterprise tools emphasize permission-aware behavior. Microsoft states Copilot “only surfaces organizational data” a user has at least view permissions for, and requires correct permission models to prevent oversharing. ²⁰ Connector setups can also be misconfigured to be “Visible to everyone,” which Microsoft warns can lead to oversharing. ²¹ Atlassian states Rovo syncs with access controls from Atlassian and connected third-party apps so users only see content they have access to. ²²

Permission-awareness doesn’t eliminate politics; it relocates it. The fight becomes: who decides the permission model, who can grant “Everyone” access, and who owns remediation of past oversharing.

Audit logs, activity histories, and compliance tooling LLM tools often create durable records. Microsoft states it stores data about Copilot interactions, including the user’s prompt and Copilot’s response (with citations), making this activity history manageable via Purview and retention policies. ²³ Google Workspace positions Gemini deployments as having “comprehensive logging and exporting capabilities for Gemini activity.” ²⁴ Glean provides admin audit logs to track administrative actions, connector changes, and configuration history. ²⁵

This introduces a surveillance-like dynamic: your knowledge work becomes inspectable. Workplace surveillance scholarship notes tensions, contradictions, and multilevel outcomes of surveillance, reinforcing that increased visibility can be simultaneously enabling (accountability) and coercive (disciplinary). ²⁶

Data governance: control, edit rights, provenance, retention

Control and edit rights When AI answers depend on what’s written, the right to edit canonical sources becomes a political resource. In federated models, domains own content but must comply with shared governance rules; in centralized models, editing is often bottlenecked by central gatekeepers, creating queue-based power (people who control the queue control prioritization). The “data mesh” framing explicitly positions decentralized ownership plus federated governance as a way to scale while avoiding siloing. ²⁷

Provenance Provenance is a technical and political requirement: it enables challenge, verification, and blame assignment. The W3C PROV data model formalizes provenance structures (entities, activities, agents) to assess quality and trustworthiness. ²⁸ RAG research highlights provenance and updateability as open issues motivating retrieval-based designs. ⁶ At the product level, Microsoft states Copilot interaction history includes citations to information used to ground the response. ²³

Provenance increases accountability but also fuels turf wars: “my doc was cited” becomes status; “your doc caused the incident” becomes liability.

Retention and deletion Retention shapes political memory. Atlassian notes Rovo retains chat/agent inputs and outputs for 30 days for safety/security purposes, and describes deletion behaviors for indexed content and connector disconnection (e.g., deleting indexed content within 30 days after disconnect). ²² Microsoft states admins can set retention policies for Copilot chat interaction data through Purview. ²³

Retention policies can be weaponized (to preserve evidence or to reduce discoverability), so governance must treat them as both compliance and power instruments. NIST’s log management guidance explicitly describes log management as generating, storing, analyzing, and disposing of log data—retention and disposal are part of the lifecycle. ²⁹

Visibility and transparency effects

AI-mediated knowledge systems increase visibility in at least four ways:

- **Discovery visibility:** content becomes findable; “hidden experts” get exposed.
- **Attribution visibility:** citations and version histories tie ideas to owners.
- **Behavioral visibility:** prompts, searches, and usage are logged. ³⁰
- **Administrative visibility:** connector changes, crawl schedules, and configuration edits are auditable. ³¹

However, visibility is rarely symmetric. Burrell distinguishes multiple forms of opacity: intentional secrecy, technical illiteracy, and opacity arising from ML scale/complexity. ³² This matters socially: those who understand the pipeline (or can claim to) gain interpretive power over what the system “really means.”

Performance metrics, incentives, and “metric gaming”

AI systems invite measurement: time saved, tickets deflected, searches-to-clicks, prompt volume, doc contribution counts, “answer acceptance” rates. This can improve accountability, but measurement regimes can also create “audit culture” pathologies where evaluation procedures become targets and reshape behavior. Strathern’s analysis of audit culture highlights how performance evaluation proliferates and changes institutions’ behavior. ³³

In AI knowledge systems, common gaming patterns include:

- Producing content optimized for ranking rather than truth.
- Over-documenting low-value artifacts to pad contribution metrics.
- “Answer theater”: generating plausible AI summaries that look productive but reduce correctness, especially outside the AI frontier (a risk highlighted by the jagged-frontier results). ³⁴

Cultural factors: knowledge sharing vs knowledge hoarding

Office politics and turf behavior can intensify around knowledge because knowledge is a source of power. Knowledge-management scholarship emphasizes that KM research often downplays power and conflict, despite their centrality. ³⁵ Empirical work on knowledge hiding distinguishes it from mere non-sharing: it is an intentional attempt to withhold requested knowledge. ³⁶ Recent qualitative research identifies **office politics as a strong organizational-level trigger for knowledge hoarding**, with employees restricting knowledge flow when sharing might jeopardize their standing. ³⁷

AI that depends on written knowledge pressures hoarders: if they do not write, their value may become less legible. Conversely, hoarders may attempt to regain advantage by controlling access, shaping narratives in docs, or resisting documentation efforts—classic turf wars in a new modality.

Legal, ethical, and risk constraints

Because jurisdiction is unspecified, the most robust approach is to assume **multi-regime compliance**:

- **Privacy principles** like purpose limitation, data minimization, accuracy, storage limitation, and integrity/confidentiality are explicit in GDPR's principles (Article 5). ³⁸
- The **EU AI Act** introduces phased obligations; the European Commission's timeline indicates major provisions apply in stages, with many rules applying by **02 Aug 2026** (including transparency rules) and earlier milestones for general provisions/AI literacy and general-purpose AI governance. ³⁹
- AI risk management frameworks (NIST AI RMF and its Generative AI Profile) provide an organizational blueprint for governing AI risks across functions and lifecycle. ⁴⁰
- Security threats (prompt injection, sensitive info disclosure, supply-chain risks) are formalized in OWASP's Top 10 for LLM Applications (v2.0/2025). ⁴¹

Vendor-specific constraints also matter. Microsoft states Copilot prompts/responses and Graph-accessed data are not used to train foundation models and that it stores interaction content and enables retention management; it also notes model updates do not change privacy/security settings. ⁴² Google emphasizes prompt injection defenses and DLP/logging controls for Gemini in Workspace. ²⁴ Atlassian notes Rovo's probabilistic outputs, permission syncing, and specific retention behaviors for chat/agents. ²²

Technical constraints that become political constraints

Technical limitations often become political flashpoints:

- **Data quality and freshness**: if AI answers are wrong, teams argue over who owns the source-of-truth and who “broke” it. Microsoft explicitly ties Copilot quality/security to data being current and governed. ⁸
- **Integration breadth**: what gets connected (tickets? chat? code? docs?) shapes which teams become visible and which remain “dark matter.”
- **Permission complexity**: misconfigured connectors can overshare. ²¹
- **Audit log ambiguity**: logs can be misinterpreted and trigger political incidents. Glean documents how Google audit logs can show “Download” events due to automated indexing via delegated access—even when no manual download occurred—and recommends validating changes before alleging unusual access. ⁴³
- **Opacity**: few people understand retrieval pipelines; technical illiteracy creates interpretive monopolies. ⁴⁴

Tactical playbooks for individuals and teams

The tactics below are organized into **overt (explicit, socially legible)** and **covert (low-visibility, plausibly deniable)** behaviors. Covert does *not* mean illegal; it means “not advertised as political.” For each tactic, the table includes effectiveness, ethical/legal risk, mitigations, and long-term career implications.

Important framing: algorithmic work tends to become a contested terrain where actors adapt strategically. ¹⁰ But the presence of audit logs, permission enforcement, and compliance regimes means harmful or unauthorized tactics are increasingly detectable and punishable, and can backfire dramatically. ⁴⁵

Comparative table of tactics vs effectiveness vs risks

Effectiveness is context-dependent; ratings assume a large software company with moderate AI adoption and typical governance.

Tactic	Overt / covert	How it creates advantage	Likely effectiveness	Key ethical/legal risks	Mitigation strategies	Long-term career implications
Become “source-of-truth” owner for a critical domain (runbooks, architecture, policy)	Overt	You control canonical artifacts the AI retrieves and cites, shaping decisions and escalations	High	Low-Medium (risk of bias/self-serving framing)	Use transparent review, include dissenting views, link to evidence; provenance-friendly authoring ⁴⁶	Strong positive if seen as steward; risk of being labeled a gatekeeper
Build a domain RAG assistant with citations and evaluation harness	Overt	Positions your team as the operational “interface” to knowledge; makes your domain legible and scalable	High	Medium (data leakage, prompt injection, overreliance)	Threat model; follow OWASP LLM Top 10; enforce permission-aware retrieval; add provenance and human review loops ⁴⁷	Strong brand as “AI-enabled platform”; also increases accountability for failures
Internal “SEO” for enterprise search/LLM retrieval (titles, headings, synonyms, Q&A structure)	Covert (benign)	Improves retrieval ranking for your artifacts; makes your framing the default answer	Medium-High	Low (unless deceptive)	Keep content accurate and balanced; avoid keyword stuffing that misleads	Positive if it improves usability; negative if discovered as manipulative

Tactic	Overt / covert	How it creates advantage	Likely effectiveness	Key ethical/legal risks	Mitigation strategies	Long-term career implications
Pre-emptive documentation before major decisions (“write the proposal into the wiki first”)	Covert (benign)	Sets agenda and constraints; AI later summarizes <i>your</i> baseline	Medium–High	Medium (agenda manipulation, exclusion)	Announce drafts early; invite review; log major changes	Can build reputation as “driver”; risk of political backlash
Join or create the AI governance forum (permissions, retention, connector scope, evaluation policy)	Overt	Governance is decision-rights; participating gives legitimate and informational power	High	Low	Anchor proposals in frameworks (AI RMF, NIST controls); keep decisions auditable ⁴⁸	High upward mobility into leadership/platform roles
Control connector scope/cadence for your domain systems	Overt–covert	Determines what is indexed and discoverable; can privilege your systems as “official”	Medium–High	Medium (oversharing, biased visibility)	Follow least privilege; avoid “Visible to everyone” patterns that overshare ⁴⁹	Platform influence; risk of being blamed for incidents
Use audit and usage analytics (legitimately) to identify demand, then author targeted docs	Covert (benign)	“Demand sensing” lets you publish what leadership/peers are searching for	Medium	Medium (privacy perception; misuse of logs)	Use aggregated/approved analytics; disclose analytics use; data minimization mindset ⁵⁰	Positive if framed as service improvement; negative if seen as surveillance

Tactic	Overt / covert	How it creates advantage	Likely effectiveness	Key ethical/legal risks	Mitigation strategies	Long-term career implications
"Permission sculpting" to preserve turf (restrict access so others must route through you)	Covert (gray)	Creates dependency; preserves gatekeeping role even with AI	Medium (short-term)	High (policy violation, discrimination, retaliation, reputational harm)	Prefer time-boxed restrictions; document rationale; align with formal classification rules; involve security/legal	Often backfires; can brand you as political and untrustworthy
Quietly seed narratives into high-authority docs (selective inclusion/exclusion of trade-offs)	Covert (gray)	AI summaries inherit your framing; dissent becomes less visible	Medium	High (misrepresentation; ethical violations)	Use explicit assumptions, trade-offs, and references; keep dissent links	Risk of long-term credibility loss if uncovered
Data poisoning / intentional misinformation in sources to manipulate AI outputs	Covert (red; not recommended)	Attempts to weaponize retrieval and summaries	Low–Medium (in secure orgs)	Extreme (security incident, termination, legal liability)	Do not do; organizations should implement provenance, review, monitoring, and access controls ⁵¹	Severe negative consequences; likely career-ending internally

Why the “lower-performer boost” matters politically

The following chart uses published effect sizes from two major field studies to illustrate a recurring pattern: AI assistance often benefits lower performers more than higher performers, changing status dynamics. ⁵²

```
xychart-beta
  title "Field evidence: AI assistance often benefits lower performers more"
  x-axis ["Support agents: avg", "Support agents: novice/low-skill", "Consultants: above avg", "Consultants: below avg"]
  y-axis "Relative improvement (%)" 0 --> 50
  bar [14, 34, 17, 43]
```

Status implications in large software companies: juniors with strong AI workflows can produce senior-like drafts faster; seniors may retain advantage primarily in frontier tasks (judgment, system design, risk calls), and by controlling governance and canonical artifacts. ⁵³

Recommended “green path” playbook for gaining influence ethically

This playbook is designed to be both effective and defensible in environments with auditability, compliance, and permission-aware AI.

Start by mapping your influence target to a control point in the knowledge supply chain:

Create a knowledge product, not just documentation. Treat your domain knowledge as a “product” with a roadmap: canonical entry points, service-level expectations for freshness, and explicit ownership. This aligns with federated governance ideas (domain ownership with shared standards). ²⁷

Design for provenance and challenge. Write in a way that can be cited and contested. Systems like Copilot explicitly store interaction records with citations, and RAG designs emphasize grounding in documents. ⁵⁴ This makes your content both more retrievable and more defensible.

Exploit “internal SEO,” but for user welfare. Use obvious titles, synonyms, FAQ headers, and decision logs so retrieval works. Algorithms adjudicate relevance; you want to be relevant by being clear. ¹⁸

Instrument outcomes, not vanity. Use accepted metrics (time-to-resolution, incident recurrence reduction, onboarding time), but beware audit-culture distortions. ³³

Make governance your ally. Join governance bodies; propose minimal, computable rules rather than manual gatekeeping. The more formal the decision rights (data governance framing), the less politics dominates. ⁵⁵

Covert but defensible tactics to protect turf without becoming a villain

Some “covert” tactics are essentially strategic communication and operational excellence:

- **Publish “decision records” immediately after key meetings,** capturing rationale, trade-offs, and links. This prevents later narrative rewrites and becomes the default retrieval target for the AI. Provenance standards emphasize capturing agents/activities involved in producing information—decision records are an organizational analog. ²⁸
- **Use permissive transparency:** share 80% broadly while protecting genuinely sensitive 20% with justified access controls, aligning with least privilege and ABAC patterns. ⁵⁶
- **Exploit audit logs as a shield, not a sword:** if a cross-team incident becomes political, use admin audit logs and interaction histories to show what changed and when—without turning it into personal surveillance. Tools like Glean explicitly position admin audit logs for troubleshooting and compliance reporting. ⁵⁷
- **Preempt “log misinterpretation scandals.”** Glean documents that automated indexing can look like mass downloads in Google audit logs and recommends validating changes before concluding unusual access. ⁴³ Knowing this lets you defuse politics that might otherwise become accusations.

Ethical and legal risk analysis and mitigation patterns

The central risk: AI turns politics into evidence

In older office politics, influence could be largely conversational and ephemeral. With AI, artifacts are:

- persisted (prompts/responses stored as interaction history),
- auditable (admin changes logged),
- sometimes exportable (for compliance),
- and increasingly regulated (AI governance and transparency obligations). ⁵⁸

This makes covert manipulation more hazardous. It also changes the ethics: “I didn’t say that” is weaker when systems store conversational history and citations. ²³

Risk categories to evaluate for any tactic

Use this checklist before choosing a tactic:

- **Privacy & surveillance risk:** Are you using logs or prompts in ways that employees would reasonably perceive as monitoring? Workplace surveillance scholarship shows surveillance can create tensions and harms even when normalized. ⁵⁹
- **Access control & oversharing risk:** Is your tactic increasing exposure of sensitive content? Microsoft explicitly warns connector permission misconfiguration (e.g., “Visible to everyone”) can overshare content. ⁶⁰
- **Security risk:** Does it create prompt injection or data leakage vulnerabilities? OWASP LLM Top 10 identifies prompt injection as a leading LLM security risk class, and Google describes prompt injection defenses for Gemini in Workspace. ⁶¹
- **Misrepresentation risk:** Are you shaping narratives in ways that would be judged deceptive if surfaced?
- **Compliance, retention, and eDiscovery risk:** Are you storing or deleting information in ways inconsistent with retention obligations? Microsoft and Atlassian both describe explicit retention behaviors for AI interactions. ⁶²
- **Regulatory exposure:** In EU contexts, AI Act timelines and transparency rules can create additional obligations and liabilities by 2026/2027. ³⁹

Mitigation strategies that preserve advantage without crossing lines

A robust mitigation pattern is to **turn influence into stewardship**:

- Prefer **transparent process power** (review workflows, published templates) over hidden manipulation—KM research argues power is inseparable from knowledge flows, so legitimacy matters. ⁶³
- Use **least privilege and ABAC** for access, with clear rationales. ABAC is explicitly intended to improve information-sharing while maintaining control. ⁶⁴
- Maintain **tamper-resistant auditability** (log management discipline) so you can defend decisions and detect misuse. ⁶⁵
- Add **provenance and citations** to stabilize trust and reduce blame-shifting. ⁶⁶

- Treat AI outputs as drafts; acknowledge the jagged frontier and require verification on critical tasks. ⁶⁷
- Align with formal risk frameworks (NIST AI RMF + Generative AI Profile) so your influence is framed as risk reduction and governance maturity, not turf capture. ⁴⁰

Long-term career implications

The “AI-mediated knowledge regime” rewards different career strategies:

- **Winners** tend to be people who can combine domain expertise with *artifact leverage*: writing canonical docs, building evaluable systems, and navigating governance forums. This maps to durable power bases (expert + informational + legitimate). ⁶⁸
- **Losers** are often those who rely on scarce tacit knowledge as a moat, particularly when AI adoption lifts junior performance and disseminates best practices faster. ⁵²
- **High-risk players** who use covert manipulation (permission sculpting, narrative suppression, data sabotage) face rising detection risk due to audit trails and retention systems. ⁴⁵

A subtle but important implication: once a team becomes the “knowledge platform,” it can become a permanent dependency—and a permanent scapegoat. To avoid the “platform blamed for everything” trap, teams should institutionalize provenance, clear decision rights, and metrics that reflect outcome quality, not just usage volume. ⁶⁹

Adoption roadmap and timeline for exploiting (and governing) the shift

This roadmap is written for a large software company but is adaptable. It includes both **organizational adoption steps** and what an individual/team can do to gain advantage ethically within each phase.

Phased roadmap

```

timeline
  title Roadmap: AI-mediated knowledge adoption and the evolving turf wars
  section Foundation
    Inventory knowledge sources, owners, ACLs : 0-30 days
    Fix oversharing / define secure defaults : 0-60 days
    Establish governance forum and decision rights : 0-60 days
  section Pilot and prove value
    Launch permission-aware enterprise search/RAG pilot : 1-3 months
    Add provenance/citations + evaluation harness : 2-4 months
    Define retention, logging, and audit workflows : 2-4 months
  section Scale across domains
    Federate content ownership + standards (templates, metadata) : 3-9 months
    Expand connectors; implement least privilege/ABAC patterns : 3-9 months
    Train AI literacy; set cultural norms for documentation : 3-9 months
  section Optimize and stabilize
    Improve ranking with feedback loops (guarded against gaming) : 6-12 months

```

Automate compliance reporting; monitor misuse patterns : 6-12 months
Rebalance incentives & career ladders for knowledge work : 9-18 months

What to do in each phase to gain influence (individual/team perspective)

Foundation phase (0–60 days): capture the map, not the territory Focus on diagnosing where power will shift. Identify (a) what will be indexed, (b) who owns it, and (c) which permissions/errors will matter. Microsoft's guidance emphasizes remediating oversharing and establishing guardrails because AI can surface content that is "overshared, ownerless, inactive, or contains sensitive data." ⁷⁰

Practical advantage: volunteer to run the inventory, own a high-value slice (one product area), and become the "person who knows where knowledge lives." This is the seed of legitimate informational power. ⁷¹

Pilot phase (1–4 months): become the team that makes AI safe and useful The winning move is to ship a narrow, high-trust pilot: one domain assistant with citations and eval. Build it like RAG research intends—retrieval + generation—so updating sources updates answers. ⁶

Practical advantage: ship an evaluation harness and a "trust story" (provenance + permissions). Vendors emphasize permission-aware access as the baseline expectation. ⁷²

Scale phase (3–9 months): fight for governance shape At scale, turf wars intensify: domains resist centralization; central platform resists fragmentation. Adopt a data-mesh-like compromise: domain ownership plus federated governance and self-serve platforms. ¹⁴

Practical advantage: become a governance "translator" who converts business conflict into computable policy (labels, access patterns, retention rules). This increases influence while reducing blame exposure. ⁷³

Optimize phase (6–18 months): control the narrative about "what works" Optimization is where metrics and politics collide. Audit culture warns that measurement systems can distort behavior. ³³ Use metrics that are hard to game (incident recurrence, cycle time reduction, onboarding time), and design controls to detect gaming (anomaly detection on edits, review policies, provenance checks). ⁷⁴

Practical advantage: own the dashboard that leadership trusts—but ensure it is grounded in defensible assumptions and privacy-respectful analytics. ⁷⁵

Organizational guardrails that reduce destructive turf wars

Organizations that want the benefits without the politics spiral should implement:

- **Clear connector permission models and least privilege** (avoid "Visible to everyone"; enforce source ACL respect). ⁷⁶
- **ABAC where roles/contexts change rapidly** (project-based access, dynamic teams). ⁷⁷
- **Auditable admin actions and incident forensics** (log management discipline; admin audit logs). ⁷⁸
- **Provenance everywhere** (citations in answers; traceability of edits; decision records). ⁷⁹

- **Security controls for LLM-specific threats** (prompt injection protections, DLP, OWASP LLM Top 10 alignment). ⁸⁰
- **AI risk management program** aligned to NIST AI RMF + GenAI profile and/or ISO/IEC 42001 where applicable. ⁸¹
- **AI literacy and transparency norms**—also consistent with EU AI Act staging (AI literacy provisions and phased obligations). ³⁹

These measures do not eliminate politics; they channel it into legitimate arenas, reducing covert conflict and lowering the risk that “AI is wrong” becomes a blame war.

Closing synthesis: the durable advantage

In AI-mediated enterprises, the durable advantage is to become the person/team that provides:

- the most trustworthy, retrievable, permission-safe artifacts,
- the best evaluation and provenance story,
- and the governance mechanisms that prevent chaos.

This strategy is ethically defensible, resilient to audits, and aligned with both technical realities (RAG dependence on written sources) and organizational realities (power embedded in knowledge flows and the contested terrain of algorithmic control). ⁸²

¹ ⁶ ⁸² Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks
https://arxiv.org/abs/2005.11401?utm_source=chatgpt.com

² ⁵² Generative AI at Work | NBER
<https://www.nber.org/papers/w31161>

³ ¹² ¹³ ³⁴ ⁵³ ⁶⁷ iab.cl
<https://www.iab.cl/wp-content/uploads/2023/11/SSRN-id4573321.pdf>

⁴ ⁷ ¹⁹ ²⁰ ²³ ³⁰ ⁴² ⁴⁵ ⁵⁰ ⁵⁴ ⁵⁸ ⁶² ⁷² ⁷⁹ Data, Privacy, and Security for Microsoft 365 Copilot | Microsoft Learn
<https://learn.microsoft.com/en-us/copilot/microsoft-365/microsoft-365-copilot-privacy>

⁵ ¹¹ ⁶⁸ ⁷¹ The Bases of Social Power
https://ia800601.us.archive.org/21/items/15341_Readings/15341_Readings/Power/French_%26_Raven_The_Bases_of_Social_Power_ch32_pp607-623_text.pdf?utm_source=chatgpt.com

⁸ ⁷⁰ Configure a secure and governed foundation for Microsoft 365 Copilot | Microsoft Learn
<https://learn.microsoft.com/en-us/copilot/microsoft-365/microsoft-365-copilot-e5-guide>

⁹ ¹⁸ 9 The Relevance of Algorithms - microsoft.com
https://www.microsoft.com/en-us/research/wp-content/uploads/2014/01/Gillespie_2014_The-Relevance-of-Algorithms.pdf?utm_source=chatgpt.com

¹⁰ Algorithms at Work: The New Contested Terrain of Control
https://journals.aom.org/doi/pdf/10.5465/annals.2018.0174?utm_source=chatgpt.com

¹⁴ ²⁷ Data Mesh Principles and Logical Architecture - Martin Fowler
https://martinfowler.com/articles/data-mesh-principles.html?utm_source=chatgpt.com

- 15 55 **Communications of the ACM - ACM Digital Library**
https://dl.acm.org/doi/fullHtml/10.1145/1629175.1629210?utm_source=chatgpt.com
- 16 63 **KNOWLEDGE MANAGEMENT : TAKING POWER DYNAMICS SERIOUSLY.**
https://warwick.ac.uk/fac/soc/wbs/conf/olkc/archive/olk4/papers/soenen.pdf?utm_source=chatgpt.com
- 17 **Untangling knowledge creation and knowledge integration in enterprise wikis**
https://link.springer.com/content/pdf/10.1007/s11573-014-0760-2.pdf?utm_source=chatgpt.com
- 21 49 60 76 **Manage access permissions for connectors - Microsoft 365 Copilot connectors | Microsoft Learn**
<https://learn.microsoft.com/en-us/microsoft-365/copilot/connectors/manage-access-permissions>
- 22 **Rovo data, privacy, and usage guidelines | Rovo | Atlassian Support**
<https://support.atlassian.com/rovo/docs/rovo-data-privacy-and-usage-guidelines/>
- 24 80 **Generative AI Security, Compliance and Privacy | Google Workspace**
<https://workspace.google.com/security/ai-privacy/>
- 25 31 57 **Audit Logs | Glean**
<https://docs.glean.com/administration/management/audit-logs/admin-audit-logs>
- 26 59 **Workplace surveillance: A systematic review, integrative framework, and ...**
https://www.sciencedirect.com/science/article/pii/S0148296323005714?utm_source=chatgpt.com
- 28 46 66 69 74 **PROV-DM: The PROV Data Model - World Wide Web Consortium (W3C)**
https://www.w3.org/TR/prov-dm/?utm_source=chatgpt.com
- 29 65 78 **SP 800-92 Rev. 1, Cybersecurity Log Management Planning Guide | CSRC**
https://csrc.nist.gov/pubs/sp/800/92/r1/ipd?utm_source=chatgpt.com
- 32 44 **How the machine 'thinks': Understanding opacity in machine learning ...**
https://journals.sagepub.com/doi/pdf/10.1177/2053951715622512?utm_source=chatgpt.com
- 33 **'Improving ratings': audi int the British University system**
https://ia803102.us.archive.org/23/items/ImprovingRatingsAuditInTheBritishUniversitySystem/%27Improving%20ratings%27-%20audit%20in%20the%20British%20University%20system.pdf?utm_source=chatgpt.com
- 35 **Knowledge Management, Power and Conflict - Springer**
https://link.springer.com/chapter/10.1007/978-3-319-71434-9_8?utm_source=chatgpt.com
- 36 **Knowledge hiding in organizations - JSTOR**
https://www.jstor.org/stable/41415737?utm_source=chatgpt.com
- 37 **The dynamics of knowledge behaviors: exploring drivers, triggers and ...**
https://www.sciencedirect.com/org/science/article/pii/S1367327025000420?utm_source=chatgpt.com
- 38 **Principles of the GDPR - European Commission**
https://commission.europa.eu/law/law-topic/data-protection/rules-business-and-organisations/principles-gdpr_en?utm_source=chatgpt.com
- 39 **Timeline for the Implementation of the EU AI Act - ai-act-service-desk ...**
https://ai-act-service-desk.ec.europa.eu/en/ai-act/timeline/timeline-implementation-eu-ai-act?utm_source=chatgpt.com
- 40 48 81 **Artificial Intelligence Risk Management Framework (AI RMF 1 - NIST**
https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf?utm_source=chatgpt.com

41 47 51 61 OWASP Top 10 for LLM Applications 2025

https://owasp.org/www-project-top-10-for-large-language-model-applications/assets/PDF/OWASP-Top-10-for-LLMs-v2025.pdf?utm_source=chatgpt.com

43 Audit Logs Behavior | Glean

<https://docs.glean.com/connectors/native/gdrive/troubleshooting/audit-logs-behavior>

56 64 73 77 SP 800-162, Guide to Attribute Based Access Control (ABAC) Definition ...

https://csrc.nist.gov/pubs/sp/800/162/upd2/final?utm_source=chatgpt.com

75 Art. 5 GDPR Principles relating to processing of personal data

https://gdpr-info.eu/art-5-gdpr/?utm_source=chatgpt.com